

DSI: AI Development Program

Part 1: Introduction to AI

University of Chicago

December 29, 2025

- Who am I
- Learning objectives for the course.
- Today: A brief introduction to the concepts around AI

Who am I

- Director Data Science Clinic
- Undergrad: UC Berkeley, PhD: UCLA
- Worked in Consulting, Video Games, AI

Some Games I've Worked on



What is an LLM?

- A **Large Language Model (LLM)** is an artificial intelligence system trained on massive amounts of text data (often trillions of words)
- These models learn patterns in language and can:
 - Generate human-like text
 - Answer questions
 - Translate languages
 - Summarize documents
 - Write code
 - And much more
- Think of an LLM as a very sophisticated autocomplete system that has read a significant portion of the internet

- A **foundation model** is a large-scale machine learning model trained on broad data (generally using self-supervision at scale) that can be adapted to a wide range of downstream tasks
- **Key characteristics:**
 - Broad training: Trained on diverse, large-scale datasets
 - General-purpose: Can be adapted to many different tasks
 - Transfer learning: Knowledge learned from training can be applied to new tasks
 - Base for specialization: Can be fine-tuned or used as-is for specific applications

- **Foundation Model** is the broader category it includes models for vision, audio, code, etc.
- **LLM (Large Language Model)** is a type of foundation model focused specifically on language
- **Examples of foundation models:**
 - **Language:** GPT-4, Claude, Llama (these are LLMs)
 - **Vision:** CLIP, DALL-E
 - **Code:** Codex, CodeLlama
 - **Multimodal:** GPT-4V (vision + language), Gemini (multimodal)

Key Terminology

- **Foundation Model:** A large-scale model trained on broad data that can be adapted to many tasks
- **LLM (Large Language Model):** A type of foundation model specifically designed for language tasks
- **Token:** The basic unit of text that an LLM processes (word, part of word, or punctuation)
- **Prompt:** The input text you give to an LLM
- **Completion/Response:** The output text generated by the LLM
- **Context Window:** The maximum number of tokens an LLM can process in a single conversation

Key Terminology (continued)

- **Temperature:** A parameter that controls randomness (Lower = more deterministic, Higher = more creative)
- **Fine-tuning:** Training an LLM on specific data to improve performance on particular tasks
- **Inference:** The process of generating text from an LLM (as opposed to training)

Context windows: what they are (and why you care)

- The **context window** is the model's working memory for a single request
- Everything the model can use must fit inside it:
 - system prompt + developer instructions
 - user messages
 - tool outputs / retrieved documents
 - the model's own previous messages (in chat)
- When you exceed the window, something gets dropped or truncated

Context windows: practical implications

- **Tradeoffs:** more context can improve grounding, but increases cost/latency
- **Recency and placement matter:** newer text is usually attended to more
- **Failure modes** to watch for:
 - forgetting earlier constraints
 - losing critical details due to truncation
 - quoting the wrong source when too many docs are included

System prompts: role and best practices

- The **system prompt** sets the highest-level behavior (tone, rules, safety, format)
- Use it to encode **non-negotiables**:
 - output format requirements (e.g., JSON only)
 - refusal and safety rules
 - grounding requirements (cite sources; do not invent facts)
- Keep it **short, explicit, testable**; avoid vague goals like “be helpful”

Context engineering (a.k.a. make the model successful)

- **Context engineering** = designing what goes into the context window so the model can reliably do the task
- Common building blocks:
 - task definition + success criteria
 - constraints and policies (what is allowed/not allowed)
 - examples (few-shot) and edge cases
 - retrieved evidence (RAG) with citations
 - tool schemas for structured interaction with code

A simple context template (recommended)

- ① **System**: rules, safety, output format
- ② **Developer**: task-specific policies (what to do, what not to do)
- ③ **User**: the request + required inputs
- ④ **Evidence**: only the relevant retrieved snippets (not entire documents)
- ⑤ **Tools**: schemas + allowed actions
- ⑥ **Final**: “Return JSON only” / “Cite sources” / “Ask clarifying questions if needed”

Example LLM Models: Commercial

- **GPT-4** (OpenAI)
 - One of the most capable models
 - Available via OpenAI API
 - Powers ChatGPT Plus
- **Claude 3.5 Sonnet** (Anthropic)
 - Strong reasoning and coding capabilities
 - Available via Anthropic API
- **Gemini Pro** (Google)
 - Google's flagship model
 - Available via Google AI Studio

Example LLM Models: Open Source

- **Hugging Face** is a platform hosting thousands of open-source models
- Popular models include:
 - **Llama 3** (Meta)
 - **Mistral** (Mistral AI)
 - **Phi** (Microsoft)
 - **Gemma** (Google)
- You can use these models via:
 - Hugging Face Transformers library (Python)
 - Hugging Face Inference API
 - Local deployment

Model Aggregators

- **Open Router** provides access to many models through a single API
 - Unified API for 100+ models
 - Compare models side-by-side
 - Pay-per-use pricing
 - Easy to switch between models

Token economics: why tokens matter

- Most LLM APIs charge by tokens:
 - **input tokens** (prompt, system prompt, retrieved docs, tool outputs)
 - **output tokens** (the model's response)
- Cost and latency generally scale with total tokens processed
- **Context engineering is cost engineering:**
 - tighter prompts → lower cost and often higher reliability
 - smaller context when possible → faster, cheaper

Pricing references: <https://openai.com/api/pricing/>

Cost math (simple model)

- Typical pricing is “\$ per 1M tokens”
- Approximate cost per call:

$$\text{cost} \approx \left(\frac{T_{in}}{10^6} \cdot p_{in} \right) + \left(\frac{T_{out}}{10^6} \cdot p_{out} \right)$$

- Engineering levers:
 - reduce retrieved context (better retrieval, smaller snippets)
 - constrain output length (max tokens, structured output)
 - pick the cheapest model that meets quality requirements
 - cache repeated context (where supported)

Pricing references: <https://openai.com/api/pricing/>

Real-world pricing examples (text)

Example (OpenAI GPT-4.1 family; per 1M tokens)

Model	Input	Output
GPT-4.1	\$2.00	\$8.00
GPT-4.1 mini	\$0.40	\$1.60
GPT-4.1 nano	\$0.10	\$0.40

- Decision implication: use smaller/cheaper models for simpler tasks (classification, routing, QC), reserve larger models for harder steps

Source: <https://openai.com/index/gpt-4-1/>

Multimodal pricing: what changes?

- Multimodal models may price **each modality differently**:
 - text tokens (input/output)
 - image generation (often priced by image or image tokens)
 - audio tokens (speech in/out) and realtime usage
 - video generation (often priced per second/resolution)
- Decision implication:
 - use multimodal only when it adds value (OCR, visual QA, audio transcription)
 - consider hybrid pipelines (cheap OCR → text-only LLM) to control cost

Pricing references: <https://openai.com/api/pricing/> and
<https://platform.openai.com/docs/pricing>

LLM Limitations

While LLMs are powerful, they have important limitations:

- ① **Knowledge Cutoff:** They only know information from their training data up to a certain date
- ② **Hallucination:** They can generate plausible-sounding but incorrect information
- ③ **No Real-World Actions:** They can't directly interact with external systems (databases, APIs, file systems)
- ④ **Context Limits:** They have maximum context window sizes
- ⑤ **Static Knowledge:** They can't learn new information after training without fine-tuning or retrieval

What is an AI Agent?

- An **AI Agent** is an LLM that can use **tools** to interact with the outside world
- While a basic LLM can only generate text based on its training data, an agent can:
 - Query databases
 - Call APIs
 - Read files
 - Execute code
 - Search the web
 - Perform actions in real-time

LLM vs Agent: Key Differences

Aspect	LLM	Agent
Capabilities	Text generation only	Text + tool usage
Knowledge	Training data only	Training data + real-time data
Actions	None	Can perform actions via tools
Interactivity	One-shot responses	Can loop and iterate

The Agent Loop

Observe → Think → Act → Observe → ...

- ① **Observe**: Receive input (user query, tool results, system state)
- ② **Think**: Process information and decide what to do next
- ③ **Act**: Execute actions (call tools, generate responses)
- ④ **Observe**: See the results and continue the loop

- **LangChain Agents**
 - Framework for building LLM applications
 - Supports multiple LLM providers
 - Tool integration and agent orchestration
- **Claude with MCP (Model Context Protocol)**
 - Claude Desktop can connect to MCP servers
 - Access custom tools and data sources
- **ChatGPT with Plugins/Code Interpreter**
- **Cursor AI**

What is a Tool?

- A **tool** is a function that an AI agent can call to interact with external systems
- Tools bridge the gap between the LLM's text generation capabilities and real-world actions
- Every tool has:
 - ① **Name:** A unique identifier (e.g., `get_player_list`)
 - ② **Description:** What the tool does and when to use it
 - ③ **Input Schema:** Parameters the tool accepts (JSON Schema format)
 - ④ **Implementation:** The actual code that executes when called

Common Tool Categories


- **Database Tools:** Query databases, insert/update/delete records
- **API Tools:** Call REST APIs, interact with web services
- **File System Tools:** Read/write files, list directories
- **Code Execution Tools:** Run Python code, execute shell commands
- **Web Tools:** Search the web, scrape websites

What is MCP?

- The **Model Context Protocol (MCP)** is an open protocol created by Anthropic
- Enables AI assistants to securely connect to external tools and data sources
- Provides a standardized way for AI assistants to discover and use capabilities from external systems

Why MCP?

- Before MCP, each AI assistant had its own way of connecting to external tools:
 - ChatGPT had plugins
 - Claude had custom integrations
 - Each system was proprietary
- MCP provides:
 - **Standardization:** One protocol works across multiple AI assistants
 - **Security:** Secure, controlled access to tools
 - **Discoverability:** AI assistants can discover available tools automatically
 - **Composability:** Tools can be combined and chained together



Key MCP Components

- ① **MCP Server:** Exposes capabilities (tools, resources, prompts) to AI assistants
- ② **MCP Client:** AI assistant that connects to servers (Claude Desktop, Cursor, etc.)
- ③ **Tools:** Functions that an AI can call to interact with external systems
- ④ **Resources:** Data sources that an AI can read (files, database tables, API endpoints)
- ⑤ **Prompts:** Pre-defined prompt templates for common tasks

How MCP Works


- ① User asks: "Get me all players"
- ② MCP Client analyzes query and selects appropriate tool
- ③ Client calls tool: `get_player_list()`
- ④ MCP Server makes HTTP request to Flask API
- ⑤ API returns player data
- ⑥ Server returns tool result to client
- ⑦ Client incorporates result into response
- ⑧ User receives: "Here are all the players: . . ."

Why not a REST API?

- MCP uses **Server-Sent Events (SSE)** for real-time communication
- **REST**: One request → wait → one response → connection closes
- **MCP/SSE**: Long-lived connection → multiple messages → streaming updates
- This enables:
 - Real-time feedback during tool execution
 - Streaming responses for long operations
 - Continuous communication between client and server

What are Claude Skills?

- **Claude Skills** (also called "Actions" or "Tool Use") is Claude's built-in capability to use tools
- When you give Claude access to tools, it can:
 - Automatically decide when to use tools
 - Call multiple tools in sequence
 - Combine tool results into coherent responses
- Claude Skills work with:
 - MCP servers (via Claude Desktop)
 - Custom API integrations
 - Function calling (via Anthropic API)



Example: Complete Interaction Flow

Scenario: User asks "What colleges do players from Washington come from?"

- ① Claude analyzes query, identifies need for player data
- ② Selects get_players tool
- ③ Calls get_players(team="WAS")
- ④ MCP Server makes GET request to Flask API
- ⑤ API queries database
- ⑥ Returns player records
- ⑦ Claude processes data, extracts unique colleges
- ⑧ Returns formatted response to user

- **LLMs** are powerful text generation systems trained on massive datasets
- **AI Agents** extend LLMs with tool-using capabilities
- **Tools** bridge the gap between LLMs and real-world systems
- **MCP** provides a standardized protocol for connecting AI assistants to external tools
- Together, these technologies enable AI systems that can interact with the real world

Questions?

Thank you!