

# Lecture 3: Improving Performance

## Context Engineering Techniques for Better Results

University of Chicago

January 25, 2026

# Outline

- 1 Where We Are
- 2 Performance Improvement
- 3 Expense Report System
- 4 Iteration 1: Breaking Down the Problem
- 5 Iteration 2: Grounding with Citations
- 6 Iteration 3: Few-Shot Examples
- 7 Other Techniques
- 8 Your Turn: Improving the Resume Scorer

# Review: What We've Covered

## Lecture 1: Foundations

- Basic AI terminology (LLMs, tokens, context windows, APIs)
- How to work with these models programmatically
- Understanding pricing and token economics

## Lecture 2: Building AI Systems

- Vertical slices - prove end-to-end flow works first
- Crawl, Walk, Run - start simple, add complexity incrementally
- Breaking problems into clear inputs and outputs

# Where We're Going

**Last Time:** We built a basic resume scoring system

- Single prompt: resume + job requirements → 0-100 score

**Today:** Learn techniques to improve AI system performance

- **Lecture:** Walk through a simple example (expense validation)
- **Work Session:** Apply these same techniques to improve the resume scorer

# Outline

- 1 Where We Are
- 2 Performance Improvement
- 3 Expense Report System
- 4 Iteration 1: Breaking Down the Problem
- 5 Iteration 2: Grounding with Citations
- 6 Iteration 3: Few-Shot Examples
- 7 Other Techniques
- 8 Your Turn: Improving the Resume Scorer

# The Core Problem

**It's difficult to know what does and does not work with AI Systems.**

- Results are inconsistent and non-deterministic
- Hard to debug when things go wrong
- Model hallucinates or makes up information
- Edge cases get mishandled
- We (often) can't explain *why* the model decided something

**This is why we use vertical slices – to identify processes that require improvement.**

# The Core Problem

## Our resume scorer from Lecture 2 had issues:

- Scores ranged from 0-100, but were they *good*?
- Results tended to be bunchy (most scores clustered in narrow ranges)
- Difficult to interpret what a "65" meant vs. a "75"
- No clear explanation for why a candidate got their score
- Hard to defend decisions to hiring managers

## How do we make this better?

- Break down what we're measuring
- Provide evidence for claims
- Make the scoring process more transparent

# The Core Problem

**Solution:** Engineering the context more carefully

- How?
  - Break complex problems into simpler steps
  - Ground the model with evidence requirements
  - Provide examples to guide behavior

# Outline

- 1 Where We Are
- 2 Performance Improvement
- 3 Expense Report System
- 4 Iteration 1: Breaking Down the Problem
- 5 Iteration 2: Grounding with Citations
- 6 Iteration 3: Few-Shot Examples
- 7 Other Techniques
- 8 Your Turn: Improving the Resume Scorer

# Expense Report Validator

**Task:** Validate employee expense reports for compliance

**Input:** Receipt/invoice text + description

**Required Outputs:**

- Expense type (Travel, Office Supplies, Meals/Entertainment)
- Date and location
- Compliance issues

**Compliance Rules:**

- ① **Amount limits:** Different limits per category
- ② **Required fields:** Travel requires trip purpose, Meals require attendees
- ③ **Prohibited expenses:** No alcohol reimbursement

# Example 1: Travel Expense

## Receipt:

Conference Badge  
AI Summit 2024  
Jan 15-17, 2024  
San Francisco, CA

Registration: \$495.00  
Processing Fee: \$15.00  

---

Total: \$510.00

## Expected Output:

- **Type:** Travel
- **Date:** Jan 15-17, 2024
- **Location:** San Francisco, CA
- **Amount:** \$510.00
- **Required fields:** Trip purpose (conference), destination
- **Compliance:** Pass

## Example 2: Meals with Alcohol

### Receipt:

The Steakhouse  
Boston, MA  
Jan 20, 2024

2x Ribeye Steak: \$68.00

1x Caesar Salad: \$14.00

1x House Wine: \$16.00

Tax: \$9.80

---

Total: \$107.80

### Expected Output:

- **Type:** Meals/Entertainment
- **Date:** Jan 20, 2024
- **Location:** Boston, MA
- **Amount:** \$107.80
- **Required fields:** Business purpose, attendees
- **Compliance:** Fail
- **Issue:** Contains prohibited alcohol (\$16.00)

# Initial Approach: Monolithic Prompt

**Initial approach:** One prompt does everything

## Prompt

Analyze this expense report and return JSON with: expense\_type, date, location, amount, compliance\_issues, required\_fields\_missing, prohibited\_items\_found.

Receipt: [receipt text here]

Check all compliance rules...

**This seems reasonable - what could go wrong?**

# Problem: Inconsistent and Unreliable

## What goes wrong:

- Misclassifies expense types (calls a conference "Office Supplies" instead of "Travel")
- Misses conditional logic (doesn't ask for trip purpose on travel expenses)
- Sometimes flags issues that don't apply to the expense type
- Results vary when you run the same receipt twice

## Why:

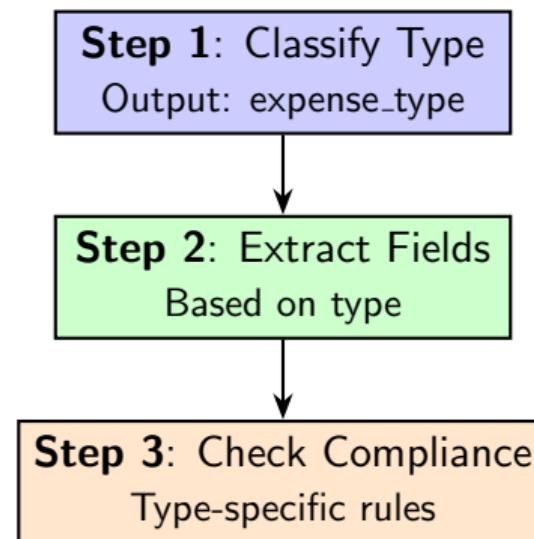
- The prompt is trying to do too many things at once
- Conditional logic ("if travel, then check X") is implicit
- Model has to keep track of multiple rules simultaneously

# Outline

- 1 Where We Are
- 2 Performance Improvement
- 3 Expense Report System
- 4 Iteration 1: Breaking Down the Problem
- 5 Iteration 2: Grounding with Citations
- 6 Iteration 3: Few-Shot Examples
- 7 Other Techniques
- 8 Your Turn: Improving the Resume Scorer

# Solution: Decompose Into Stages

**Better approach:** Break into sequential steps with clear outputs



**Key idea:** Each step is simple and has a single responsibility

# Step 1: Classify Expense Type

**First prompt:** Just classify the type

## Classification Prompt

Classify this expense into one of three categories:

- Travel (flights, hotels, conferences, transportation)
- Office Supplies (equipment, software, books)
- Meals/Entertainment (restaurants, client dinners)

Receipt: [Dinner at restaurant, \$48.50]

Return JSON: {"expense\_type": "..."}  
{"expense\_type": "Meals/Entertainment"}

**Result:** {"expense\_type": "Meals/Entertainment"}

## Step 2: Conditional Logic in Code

Now we can apply type-specific logic in our code:

### Python Logic

```
if expense_type == "Travel":  
    # Extract: trip_purpose, destination, dates  
    run_travel_prompt()  
  
elif expense_type == "Meals/Entertainment":  
    # Extract: attendees, business_purpose  
    run_meals_ent_prompt()  
  
elif expense_type == "Office Supplies":  
    run_office_supplies_prompt()
```

**Key insight:** LLMs are great at classification; code is great at conditional logic

# Results After Decomposition

## What improved:

- Classification is more consistent and reliable
- Type-specific rules are explicit in code
- Easy to debug (we can inspect the output of each step)
- Can test classification independently from validation

## Takeaway:

**Simplify each prompt by breaking complex tasks into stages**

Use LLMs for unstructured→structured transformation

Use code for logic and conditional rules

# Outline

- 1 Where We Are
- 2 Performance Improvement
- 3 Expense Report System
- 4 Iteration 1: Breaking Down the Problem
- 5 Iteration 2: Grounding with Citations
- 6 Iteration 3: Few-Shot Examples
- 7 Other Techniques
- 8 Your Turn: Improving the Resume Scorer

## New Problem: Hallucination

**Context:** We're checking for prohibited expenses (alcohol)

**Prompt:** Does this receipt contain alcohol? Return JSON: {"contains\_alcohol": true/false, "reason": "..."}  
Receipt: [Restaurant bill, \$48.50, two entrees and soft drinks]

**Result:** {"contains\_alcohol": true, "reason": "Likely includes wine or beer with dinner"}

**Problem:** The model *assumed* there was alcohol when the receipt never mentioned it!

# Why This Happens

**LLMs are pattern matchers**, not database lookups:

- Restaurant + dinner → often includes alcohol in training data
- Model generates "plausible" answers based on patterns
- Without explicit evidence requirement, it fills in gaps

**This is dangerous:**

- False rejections (denying valid expenses)
- Compliance issues (wrong reasons in audit trail)
- Loss of trust in the system

# Solution: Require Citations

**Better prompt:** Demand evidence from the receipt or ask to verify before returning (which can add \$\$\$)

## Grounded Prompt

Does this receipt contain alcohol?

**You MUST cite exact quotes from the receipt to support your answer.**

If alcohol is present, provide the exact line items that mention it.

Receipt: [Restaurant bill, \$48.50, two entrees and soft drinks]

Return JSON:

```
{  
  "contains_alcohol": true/false,  
  "evidence": ["exact quote 1", "exact quote 2"],  
  "reason": "explanation based on evidence"  
}
```

# Results After Grounding

## New result:

```
{  
    "contains_alcohol": false,  
    "evidence": [],  
    "reason": "No alcohol items found in receipt.  
                Items listed: 2x Lunch Special, 2x Soft Drink"  
}
```

## What improved:

- Model is forced to ground claims in actual receipt text
- Hallucinations are reduced significantly
- We get an audit trail (can verify the evidence ourselves)
- False positives drop dramatically

# Grounding Best Practices

## How to apply grounding:

- ① Require exact quotes or citations for any factual claims
- ② Ask for evidence *before* conclusions
- ③ Validate that evidence exists in the source material
- ④ Use structured output with separate "evidence" field

## When to use it:

- Compliance/regulatory contexts (must justify decisions)
- Extracting facts from documents
- Any scenario where hallucination is risky

**Citations turn the LLM from a "creative writer" into an "evidence-based analyst"**

# Outline

- 1 Where We Are
- 2 Performance Improvement
- 3 Expense Report System
- 4 Iteration 1: Breaking Down the Problem
- 5 Iteration 2: Grounding with Citations
- 6 Iteration 3: Few-Shot Examples
- 7 Other Techniques
- 8 Your Turn: Improving the Resume Scorer

# Another Problem: Edge Cases

**Context:** Classification works well for obvious cases

**But what about ambiguous expenses?**

- Conference registration: Travel or Office Supplies?
- Team building dinner: Meals or Office Supplies?
- Online course subscription: Office Supplies or ???
- Software for a specific project: Office Supplies or ???

**Problem:** Model makes inconsistent decisions on edge cases

- Conference registration sometimes classified as "Office Supplies"
- No clear policy guidance in the prompt

# Solution: Provide Examples

**Better prompt:** Show the model what good classification looks like

## Few-Shot Prompt

Classify this expense. Here are some examples:

**Example 1:** Receipt: Conference registration, AI Summit 2024, \$500 Classification: Travel  
(conferences are travel-related)

**Example 2:** Receipt: Office chair from Staples, \$200 Classification: Office Supplies (furniture and equipment)

**Example 3:** Receipt: Team dinner at local restaurant, \$120 Classification: Meals/Entertainment  
(company-sponsored meals)

Now classify: Receipt: [Your receipt here]

# Types of Examples to Include

**Good few-shot examples include:**

- ① **Positive examples:** Correct classifications
- ② **Edge cases:** Ambiguous situations with your preferred handling
- ③ **Negative examples:** Common mistakes to avoid

**Example - Negative case:** Receipt: Laptop purchased for work, \$1,200

*Incorrect:* Meals/Entertainment

*Correct:* Office Supplies (technology equipment)

*Reason:* Even though expensive, hardware is office supplies, not travel

**How many examples?** 2-5 is usually enough; more isn't always better

# Results After Few-Shot Examples

## What improved:

- Edge cases are now handled consistently
- Model learns your organization's specific policies
- Classification accuracy improves on ambiguous cases
- Fewer surprises and unexpected categorizations

## Tradeoff:

- More tokens in prompt (costs more, takes longer)
- Need to maintain example set
- But: much more reliable results

**Examples teach the model your specific policies and edge case handling**

# Outline

- 1 Where We Are
- 2 Performance Improvement
- 3 Expense Report System
- 4 Iteration 1: Breaking Down the Problem
- 5 Iteration 2: Grounding with Citations
- 6 Iteration 3: Few-Shot Examples
- 7 Other Techniques
- 8 Your Turn: Improving the Resume Scorer

# Additional Performance Techniques

Beyond decomposition, grounding, and examples:

## ① Role/Persona Specification

- "You are a careful compliance officer reviewing expenses..."
- Sets tone and behavior expectations

## ② Providing Domain Context

- Include relevant company policies in the prompt
- Specific rules and thresholds from your domain

## ③ Adding Guardrails

- Structured output schemas (JSON with required fields)
- Validation logic to catch impossible outputs
- Allowed/blocked lists (e.g., only these 3 expense types)

# More Advanced Techniques

## ④ Chain-of-Thought Prompting

- Ask model to "show your reasoning step by step"
- Improves accuracy on complex logic problems
- Makes debugging easier

## ⑤ Temperature Tuning

- Lower temperature (0.0-0.3) for consistency
- Higher temperature (0.7-1.0) for creativity
- Classification and extraction: use low temperature

## ⑥ Output Validation

- Check LLM outputs with deterministic code
- Re-prompt if validation fails
- Combine AI flexibility with code reliability

# The Three Core Techniques (Summary)

## 1. Decompose

Break complex prompts into stages  
Use code for conditional logic

## 2. Ground

Require citations  
Demand evidence  
Reduce hallucination

## 3. Examples

Few-shot prompting  
Show edge cases  
Guide behavior

**Use these in combination for best results**

# Outline

- 1 Where We Are
- 2 Performance Improvement
- 3 Expense Report System
- 4 Iteration 1: Breaking Down the Problem
- 5 Iteration 2: Grounding with Citations
- 6 Iteration 3: Few-Shot Examples
- 7 Other Techniques
- 8 Your Turn: Improving the Resume Scorer

# Today's Workflow

You will work with 3 randomly sampled resumes:

① **Baseline:** Run the monolithic scorer

- Single prompt: resume → 0-100 score
- Collect results in a DataFrame

② **TODO 1:** Extract years of experience

- Write a focused prompt to extract only years of experience
- Require citations/evidence from resume
- Run on all 3 samples, create DataFrame

③ **TODO 2:** Extract and score technologies

- Extract technologies/skills from resume
- Compare against required technologies in job req
- Score 0-100 based on match percentage

# Today's Workflow (continued)

## ④ TODO 3: Combine and compare

- Merge year data with technology scores
- Create a composite score from extracted features
- Compare against monolithic 0-100 scores
- Analyze: Which is more consistent? Which explains better?

## ⑤ TODO 4: Extensions (pick one or more)

- Education requirements extraction
- Leadership/mentoring experience
- Scale to 20 resumes and compare costs

# Key Takeaways for Your Work

## Apply the three core techniques:

- **Decomposition:** Break scoring into feature extraction steps
- **Grounding:** Require citations for years of experience
- **Examples:** (Optional) Add few-shot examples for edge cases

## What to look for:

- Are decomposed scores more explainable than monolithic scores?
- Can you justify each score with specific extracted features?
- How do token costs compare between approaches?
- Which approach would you trust more in production?

**Remember:** The goal is *transparency* and *consistency*, not just higher scores!

# Let's Get Started!

Open the notebook:

`lecture_3_resume_scorer_improvement.ipynb`

Work through TODOs 1-4

Compare your results

Discuss with your team